

MLproject1_benz_shibna

February 2, 2023

1 Importing important Libraries

```
[3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
[4]: train_data=pd.read_csv("train.csv")
test_data=pd.read_csv("test.csv")
```

1.0.1 Checking the no:of rows and columns each file has

```
[5]: (train_data.shape , test_data.shape )
```

```
[5]: ((4209, 378), (4209, 377))
```

```
[6]: train_data.head(10)
```

```
[6]:
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	\
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	
5	18	92.93	t	b	e	c	d	g	h	s	...	0	0	1	0	0	
6	24	128.76	al	r	e	f	d	f	h	s	...	0	0	0	0	0	
7	25	91.91	o	l	as	f	d	f	j	a	...	0	0	0	0	0	
8	27	108.67	w	s	as	e	d	f	i	h	...	1	0	0	0	0	
9	30	126.99	j	b	aq	c	d	f	a	e	...	0	0	1	0	0	

	X380	X382	X383	X384	X385
0	0	0	0	0	0
1	0	0	0	0	0
2	0	1	0	0	0
3	0	0	0	0	0

4	0	0	0	0	0
5	0	0	0	0	0
6	0	0	0	0	0
7	0	0	0	0	0
8	0	0	0	0	0
9	0	0	0	0	0

[10 rows x 378 columns]

```
[7]: test_data.head()
```

```
[7]:   ID  X0 X1  X2 X3 X4 X5 X6 X8  X10  ...  X375  X376  X377  X378  X379  X380  \
0    1  az  v   n  f  d  t  a  w    0  ...    0    0    0    1    0    0
1    2   t  b  ai  a  d  b  g  y    0  ...    0    0    1    0    0    0
2    3  az  v  as  f  d  a  j  j    0  ...    0    0    0    1    0    0
3    4  az  l   n  f  d  z  l  n    0  ...    0    0    0    1    0    0
4    5   w  s  as  c  d  y  i  m    0  ...    1    0    0    0    0    0
```

	X382	X383	X384	X385
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 377 columns]

```
[8]: train_data.columns
```

```
[8]: Index(['ID', 'y', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8',
...
        'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',
        'X385'],
        dtype='object', length=378)
```

```
[9]: test_data.columns
```

```
[9]: Index(['ID', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8', 'X10',
...
        'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',
        'X385'],
        dtype='object', length=377)
```

1.0.2 Missing value

```
[10]: missing_data = train_data.isnull().sum(axis=0).reset_index()
print(missing_data.info())
missing_data.columns = ['count_name', 'missing_count']
missing_data = missing_data.loc[missing_data['missing_count'] > 0]
missing_data = missing_data.sort_values(by='missing_count')
missing_data
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 378 entries, 0 to 377
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   index    378 non-null    object
1   0         378 non-null    int64
dtypes: int64(1), object(1)
memory usage: 6.0+ KB
None
```

```
[10]: Empty DataFrame
Columns: [count_name, missing_count]
Index: []
```

No missing value

```
[11]: missing_data = test_data.isnull().sum(axis=0).reset_index()
print(missing_data.info())
missing_data.columns = ['count_name', 'missing_count']
missing_data = missing_data.loc[missing_data['missing_count'] > 0]
missing_data = missing_data.sort_values(by='missing_count')
missing_data
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 377 entries, 0 to 376
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   index    377 non-null    object
1   0         377 non-null    int64
dtypes: int64(1), object(1)
memory usage: 6.0+ KB
None
```

```
[11]: Empty DataFrame
Columns: [count_name, missing_count]
Index: []
```

no missing data in test data also

[]:

1.0.3 Checking the variance in data_train set

[12]: `train_filter_var=np.var(train_data)`

```
C:\Users\shibn\anaconda3\lib\site-packages\numpy\core\fromnumeric.py:3721:
FutureWarning: Dropping of nuisance columns in DataFrame reductions (with
'numeric_only=None') is deprecated; in a future version this will raise
TypeError. Select only valid columns before calling the reduction.
    return var(axis=axis, dtype=dtype, out=out, ddof=ddof, **kwargs)
```

[13]: `train_filter_var`

```
[13]: ID      5.940524e+06
      y      1.607285e+02
      X10    1.312780e-02
      X11    0.000000e+00
      X12    6.944063e-02
      ...
      X380    8.012675e-03
      X382    7.544954e-03
      X383    1.660337e-03
      X384    4.749465e-04
      X385    1.423485e-03
      Length: 370, dtype: float64
```

[14]: `print(train_filter_var==0)`

```
ID      False
y      False
X10     False
X11      True
X12     False
      ...
X380    False
X382    False
X383    False
X384    False
X385    False
      Length: 370, dtype: bool
```

[15]: `zero_variance_train=train_filter_var[train_filter_var==0]`

[16]: `zero_variance_train.shape`

```
[16]: (12,)
```

```
[17]: print(zero_variance_train)
```

```
X11      0.0
X93      0.0
X107     0.0
X233     0.0
X235     0.0
X268     0.0
X289     0.0
X290     0.0
X293     0.0
X297     0.0
X330     0.0
X347     0.0
dtype: float64
```

```
[18]: zero_variance_train.keys()
```

```
[18]: Index(['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293',
          'X297', 'X330', 'X347'],
          dtype='object')
```

```
[19]: zero_var_keydata=zero_variance_train.keys()
```

1.0.4 Eliminate the columns which has zero variance

```
[20]: new_train_data=train_data.drop(zero_var_keydata ,axis=1)
```

```
[21]: new_train_data.shape
```

```
[21]: (4209, 366)
```

1.0.5 to check which all columns are dtype==object

```
[22]: new_train_data.dtypes==object
```

```
[22]: ID      False
      y      False
      X0      True
      X1      True
      X2      True
      ...
      X380    False
```

```

X382    False
X383    False
X384    False
X385    False
Length: 366, dtype: bool

```

```
[23]: x=(new_train_data.dtypes==object)
```

```
[24]: x.value_counts()
```

```

[24]: False    358
      True      8
      dtype: int64

```

```
[25]: object_train_cols = list(x[x].index)
```

```
[26]: print(object_train_cols)
```

```
['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']
```

```
[ ]:
```

1.0.6 Checking the variance in data_test set

```
[27]: ### using sklearn
```

```

[28]: #from sklearn.feature_selection import VarianceThreshold
      #selector = VarianceThreshold(threshold=0)
      #selector.fit_transform(test_data)

```

```
[29]: test_filter_var=np.var(test_data)
```

```

C:\Users\shibn\anaconda3\lib\site-packages\numpy\core\fromnumeric.py:3721:
FutureWarning: Dropping of nuisance columns in DataFrame reductions (with
'numeric_only=None') is deprecated; in a future version this will raise
TypeError. Select only valid columns before calling the reduction.
    return var(axis=axis, dtype=dtype, out=out, ddof=ddof, **kwargs)

```

```
[30]: test_filter_var
```

```

[30]: ID          5.869917e+06
      X10         1.864563e-02
      X11         2.375297e-04
      X12         6.883438e-02
      X13         5.733136e-02
      ...

```

```
X380    8.012675e-03
X382    8.713410e-03
X383    4.749465e-04
X384    7.122504e-04
X385    1.660337e-03
Length: 369, dtype: float64
```

```
[31]: print(test_filter_var==0)
```

```
ID      False
X10      False
X11      False
X12      False
X13      False
...
X380     False
X382     False
X383     False
X384     False
X385     False
Length: 369, dtype: bool
```

```
[32]: zero_test_var=test_filter_var[test_filter_var==0]
```

```
[33]: zero_test_var.shape
```

```
[33]: (5,)
```

```
[34]: zero_test_var.keys()
```

```
[34]: Index(['X257', 'X258', 'X295', 'X296', 'X369'], dtype='object')
```

```
[35]: zero_var_cols=zero_test_var.keys()
```

```
[36]: new_test_data=test_data.drop(zero_var_cols,axis=1)
```

```
[37]: new_test_data.shape
```

```
[37]: (4209, 372)
```

1.0.7 to check which all columns are object data type

```
[38]: x=(new_test_data.dtypes==object)
```

```
[39]: x.value_counts()
```

```
[39]: False    364
      True      8
      dtype: int64
```

```
[40]: object_test_cols = list(x[x].index)
```

```
[41]: print(object_test_cols)
```

```
['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']
```

2 Applying label encoding

```
[42]: from sklearn.preprocessing import LabelEncoder
      label_Encoder=LabelEncoder()
```

```
[43]: for i in object_train_cols:
      new_train_data[i]=label_Encoder.fit_transform(new_train_data[i])
      #new_data_test[i]=label_Encoder.fit_transform(new_test_data[i])
```

```
[44]: t=new_train_data.dtypes==object
```

```
[45]: t.value_counts()
```

```
[45]: False    366
      dtype: int64
```

The above step is to confirm that the label encoder have converted everything to integer that is why no true value is found as per the condition given any datatypes is equal to object

```
[46]: new_train_data.head()
```

```
[46]:
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	\
0	0	130.81	32	23	17	0	3	24	9	14	...	0	0	1	0	
1	6	88.53	32	21	19	4	3	28	11	14	...	1	0	0	0	
2	7	76.26	20	24	34	2	3	27	9	23	...	0	0	0	0	
3	9	80.62	20	21	34	5	3	27	11	4	...	0	0	0	0	
4	13	78.02	20	23	34	5	3	12	3	13	...	0	0	0	0	

	X379	X380	X382	X383	X384	X385
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	1	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0

[5 rows x 366 columns]

```
[47]: new_train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Columns: 366 entries, ID to X385
dtypes: float64(1), int32(8), int64(357)
memory usage: 11.6 MB
```

from the above, we came to know that the data is converted to numerical

Applying label encoding to test data

```
[48]: for i in object_test_cols:
      #new_train_data[i]=label_Encoder.fit_transform(new_train_data[i])
      new_test_data[i]=label_Encoder.fit_transform(new_test_data[i])
```

```
[49]: t=new_test_data.dtypes==object
```

```
[50]: t.value_counts()
```

```
[50]: False      372
      dtype: int64
```

The above step is to confirm that the label encoder have converted everything to integer that vis why no true value is found as per the condition given any datatypes is equal to object

```
[51]: new_test_data.head()
```

```
[51]:
```

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	\
0	1	21	23	34	5	3	26	0	22	0	...	0	0	0	1	0	
1	2	42	3	8	0	3	9	6	24	0	...	0	0	1	0	0	
2	3	21	23	17	5	3	0	9	9	0	...	0	0	0	1	0	
3	4	21	13	34	5	3	31	11	13	0	...	0	0	0	1	0	
4	5	45	20	17	2	3	30	8	12	0	...	1	0	0	0	0	

	X380	X382	X383	X384	X385
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

[5 rows x 372 columns]

```
[52]: new_test_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 4209 entries, 0 to 4208  
Columns: 372 entries, ID to X385  
dtypes: int32(8), int64(364)  
memory usage: 11.8 MB
```

from the above, we came to know that the data is converted to numerical

2.0.1 last and final stage to split as feature and target

```
[53]: train_data_Xfeatures=new_train_data.drop(["ID","y"],axis=1)  
train_data_ytarget=new_train_data['y']
```

```
[71]: idOf_test=new_test_data['ID'].values  
test_data_Xfeatures=new_test_data.drop(["ID"],axis=1)
```

3 Perform dimensionality reduction

```
[73]: # Linear dimensionality reduction using Singular Value Decomposition of  
# the data to project it to a lower dimensional space.  
from sklearn.decomposition import PCA  
pca=PCA(n_components=12,random_state=46)  
train_data_Xfeatures=pca.fit_transform(train_data_Xfeatures)  
test_data_Xfeatures=pca.fit_transform(test_data_Xfeatures)
```

4 Training using XGBoost

```
[74]: import xgboost as xgb  
from xgboost.sklearn import XGBRegressor  
import datetime  
from sklearn.model_selection import GridSearchCV  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import r2_score
```

```
[75]: X_train,X_test,y_train,y_test=train_test_split(train_data_Xfeatures,train_data_ytarget,test_si  
↪ 3,random_state=42)
```

```
[76]: print(X_train.shape,y_train.shape)
```

```
(2946, 12) (2946,)
```

```
[77]: print(X_test.shape,y_test.shape)
```

```
(1263, 12) (1263,)
```

```
[78]: xgb1 = XGBRegressor()  
parameters = {'nthread':[4], #when use hyperthread, xgboost may become slower  
              'objective':['reg:linear'],  
              'learning_rate': [0.1,0.2,.03, 0.05, .07], #so called `eta` value  
              'max_depth': [1,2,3,4,5, 6, 7],  
              'min_child_weight': [4],  
              'silent': [1],  
              'subsample': [0.7],  
              'colsample_bytree': [0.7],  
              'n_estimators': [500]}
```

```
[79]: xgb_grid = GridSearchCV(xgb1,  
                             parameters,  
                             cv = 3,  
                             n_jobs = 4,  
                             verbose=True)  
  
xgb_grid.fit(X_train,  
            y_train)
```

Fitting 3 folds for each of 35 candidates, totalling 105 fits

[21:32:30] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-03de431ba26204c4d-1/xgboost/xgboost-ci-windows/src/objective/regression_obj.cu:213: reg:linear is now deprecated in favor of reg:squarederror.

[21:32:30] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-03de431ba26204c4d-1/xgboost/xgboost-ci-windows/src/learner.cc:767: Parameters: { "silent" } are not used.

```
[79]: GridSearchCV(cv=3,  
                  estimator=XGBRegressor(base_score=None, booster=None,  
                                         callbacks=None, colsample_bylevel=None,  
                                         colsample_bynode=None,  
                                         colsample_bytree=None,  
                                         early_stopping_rounds=None,  
                                         enable_categorical=False, eval_metric=None,  
                                         feature_types=None, gamma=None, gpu_id=None,  
                                         grow_policy=None, importance_type=None,  
                                         interaction_constraints=None,  
                                         learning_rate=None, m...  
                                         monotone_constraints=None, n_estimators=100,  
                                         n_jobs=None, num_parallel_tree=None,
```

```

        predictor=None, random_state=None, ...),
    n_jobs=4,
    param_grid={'colsample_bytree': [0.7],
                'learning_rate': [0.1, 0.2, 0.03, 0.05, 0.07],
                'max_depth': [1, 2, 3, 4, 5, 6, 7],
                'min_child_weight': [4], 'n_estimators': [500],
                'nthread': [4], 'objective': ['reg:linear'],
                'silent': [1], 'subsample': [0.7]},
    verbose=True)

```

```

[80]: print(xgb_grid.best_score_)
      print(xgb_grid.best_params_)

```

```

0.4693980501828196
{'colsample_bytree': 0.7, 'learning_rate': 0.03, 'max_depth': 4,
 'min_child_weight': 4, 'n_estimators': 500, 'nthread': 4, 'objective':
 'reg:linear', 'silent': 1, 'subsample': 0.7}

```

```

[81]: #data_dmatrix = xgb.DMatrix(data=X,label=y)

```

```

[82]: #from xgboost.sklearn import XGBRegressor
      #from xgboost import XGBRegressor
      #xg_reg = xgb1(objective="reg:linear", learning_rate = 0.03, max_depth = 4,
      ↪min_child_weight = 4, n_estimators = 500, subsample = 0.7)

```

4.0.1 Training using XGBoost

```

[83]: d_train = xgb.DMatrix(X_train, label=y_train)
      d_valid = xgb.DMatrix(X_test, label=y_test)
      d_test = xgb.DMatrix(test_data_Xfeatures)

```

```

[84]: params={}
      params['objective'] = 'reg:linear'
      params['eta'] = 0.02
      params['max_depth'] = 4
      def xgb_r2_score(preds, dtrain):
          labels = dtrain.get_label()
          return 'r2', r2_score(labels, preds)
      watchlist = [(d_train, 'train'), (d_valid, 'valid')]
      clf = xgb.train(params, d_train, 1000, watchlist, early_stopping_rounds=50,
      ↪feval=xgb_r2_score, maximize=True, verbose_eval=10)

```

```

[21:32:32] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-
group-i-03de431ba26204c4d-1/xgboost/xgboost-ci-
windows/src/objective/regression_obj.cu:213: reg:linear is now deprecated in
favor of reg:squarederror.

```

[0]	train-rmse:98.73858	train-r2:-62.78156	valid-rmse:99.53598
	valid-r2:-54.41767		
[10]	train-rmse:80.91897	train-r2:-41.83730	valid-rmse:81.73402
	valid-r2:-36.36748		
[20]	train-rmse:66.39288	train-r2:-27.83795	valid-rmse:67.24898
	valid-r2:-24.29644		

C:\Users\shibn\anaconda3\lib\site-packages\xgboost\core.py:617: FutureWarning:
Pass `evals` as keyword args.

warnings.warn(msg, FutureWarning)

C:\Users\shibn\anaconda3\lib\site-packages\xgboost\training.py:39: UserWarning:
`feval` is deprecated, use `custom_metric` instead. They have different
behavior when custom objective is also used. See
https://xgboost.readthedocs.io/en/latest/tutorials/custom_metric_obj.html for
details on the `custom_metric`.

warnings.warn(

[30]	train-rmse:54.56281	train-r2:-18.47668	valid-rmse:55.47389
	valid-r2:-16.21335		
[40]	train-rmse:44.94648	train-r2:-12.21639	valid-rmse:45.91759
	valid-r2:-10.79359		
[50]	train-rmse:37.15140	train-r2:-8.02967	valid-rmse:38.19713
	valid-r2:-7.16111		
[60]	train-rmse:30.85581	train-r2:-5.22867	valid-rmse:31.98680
	valid-r2:-4.72308		
[70]	train-rmse:25.78684	train-r2:-3.35029	valid-rmse:27.01631
	valid-r2:-3.08263		
[80]	train-rmse:21.72926	train-r2:-2.08895	valid-rmse:23.06947
	valid-r2:-1.97689		
[90]	train-rmse:18.50854	train-r2:-1.24112	valid-rmse:19.97030
	valid-r2:-1.23078		
[100]	train-rmse:15.97889	train-r2:-0.67038	valid-rmse:17.56292
	valid-r2:-0.72537		
[110]	train-rmse:14.01209	train-r2:-0.28448	valid-rmse:15.71586
	valid-r2:-0.38154		
[120]	train-rmse:12.50901	train-r2:-0.02369	valid-rmse:14.33608
	valid-r2:-0.14961		
[130]	train-rmse:11.36021	train-r2:0.15571	valid-rmse:13.31214
	valid-r2:0.00875		
[140]	train-rmse:10.50934	train-r2:0.27744	valid-rmse:12.56300
	valid-r2:0.11718		
[150]	train-rmse:9.87095	train-r2:0.36256	valid-rmse:12.02727
	valid-r2:0.19086		
[160]	train-rmse:9.39646	train-r2:0.42237	valid-rmse:11.63483
	valid-r2:0.24280		
[170]	train-rmse:9.04274	train-r2:0.46504	valid-rmse:11.35249
	valid-r2:0.27911		
[180]	train-rmse:8.78203	train-r2:0.49544	valid-rmse:11.15587
	valid-r2:0.30386		

[190]	train-rmse:8.58170 valid-r2:0.32124	train-r2:0.51820	valid-rmse:11.01577
[200]	train-rmse:8.44132 valid-r2:0.33289	train-r2:0.53383	valid-rmse:10.92084
[210]	train-rmse:8.33111 valid-r2:0.34192	train-r2:0.54592	valid-rmse:10.84662
[220]	train-rmse:8.23994 valid-r2:0.34829	train-r2:0.55581	valid-rmse:10.79402
[230]	train-rmse:8.15909 valid-r2:0.35297	train-r2:0.56448	valid-rmse:10.75521
[240]	train-rmse:8.09280 valid-r2:0.35570	train-r2:0.57153	valid-rmse:10.73246
[250]	train-rmse:8.03866 valid-r2:0.35736	train-r2:0.57725	valid-rmse:10.71860
[260]	train-rmse:7.98870 valid-r2:0.35916	train-r2:0.58248	valid-rmse:10.70360
[270]	train-rmse:7.95001 valid-r2:0.35990	train-r2:0.58652	valid-rmse:10.69745
[280]	train-rmse:7.91634 valid-r2:0.36071	train-r2:0.59001	valid-rmse:10.69069
[290]	train-rmse:7.88756 valid-r2:0.36170	train-r2:0.59299	valid-rmse:10.68239
[300]	train-rmse:7.84931 valid-r2:0.36223	train-r2:0.59693	valid-rmse:10.67797
[310]	train-rmse:7.81953 valid-r2:0.36240	train-r2:0.59998	valid-rmse:10.67649
[320]	train-rmse:7.78954 valid-r2:0.36247	train-r2:0.60304	valid-rmse:10.67590
[330]	train-rmse:7.76483 valid-r2:0.36217	train-r2:0.60556	valid-rmse:10.67846
[340]	train-rmse:7.73786 valid-r2:0.36257	train-r2:0.60829	valid-rmse:10.67514
[350]	train-rmse:7.71039 valid-r2:0.36214	train-r2:0.61107	valid-rmse:10.67867
[360]	train-rmse:7.67913 valid-r2:0.36252	train-r2:0.61421	valid-rmse:10.67556
[370]	train-rmse:7.64938 valid-r2:0.36244	train-r2:0.61720	valid-rmse:10.67616
[380]	train-rmse:7.61585 valid-r2:0.36244	train-r2:0.62055	valid-rmse:10.67623
[390]	train-rmse:7.58177 valid-r2:0.36262	train-r2:0.62394	valid-rmse:10.67473
[400]	train-rmse:7.55057 valid-r2:0.36276	train-r2:0.62702	valid-rmse:10.67351
[410]	train-rmse:7.52849 valid-r2:0.36244	train-r2:0.62920	valid-rmse:10.67617
[420]	train-rmse:7.50319 valid-r2:0.36237	train-r2:0.63169	valid-rmse:10.67677

```

[430]   train-rmse:7.47070      train-r2:0.63487      valid-rmse:10.67192
valid-r2:0.36295
[440]   train-rmse:7.44340      train-r2:0.63754      valid-rmse:10.66981
valid-r2:0.36320
[450]   train-rmse:7.42140      train-r2:0.63968      valid-rmse:10.67071
valid-r2:0.36310
[460]   train-rmse:7.39629      train-r2:0.64211      valid-rmse:10.66964
valid-r2:0.36322
[470]   train-rmse:7.37310      train-r2:0.64435      valid-rmse:10.66897
valid-r2:0.36330
[480]   train-rmse:7.34658      train-r2:0.64690      valid-rmse:10.66844
valid-r2:0.36337
[490]   train-rmse:7.32267      train-r2:0.64920      valid-rmse:10.66970
valid-r2:0.36322
[500]   train-rmse:7.30167      train-r2:0.65121      valid-rmse:10.66919
valid-r2:0.36328
[510]   train-rmse:7.27660      train-r2:0.65360      valid-rmse:10.67193
valid-r2:0.36295
[520]   train-rmse:7.25104      train-r2:0.65603      valid-rmse:10.67042
valid-r2:0.36313
[526]   train-rmse:7.23652      train-r2:0.65741      valid-rmse:10.67395
valid-r2:0.36271

```

4.0.2 Predict the test data values using XGBoost

```

[85]: pred_test=clf.predict(d_test)
      prediction_data=pd.DataFrame()
      prediction_data['ID']=idOf_test
      prediction_data['y']=pred_test
      prediction_data.to_csv('predXGB.csv', index=False)
      prediction_data.head()

```

```

[85]:   ID      y
0    1  77.342758
1    2  97.923904
2    3  83.637695
3    4  78.497604
4    5 110.838745

```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```

[ ]:
[ ]:
[ ]:
[ ]:
[ ]:
[ ]:
[ ]:
[ ]:
[ ]:
[86]: #params = {'n_estimators':[10, 20, 40, 80], 'max_depth':[1,2,3,6,10],
        ↳ 'learning_rate' : [0.1, 0.2, 0.3, 0.5], 'min_child_weight' : [1, 2, 3, 4, 5],
        ↳ 'subsample' : [0.5, 0.6, 0.7, 0.8, 1.0]}
        #grid_search = GridSearchCV(xgb_clf, params, cv = 3, n_jobs = -1)
        #grid_search.fit(X_train, y_train)

[87]: #grid_search.best_params_

[ ]: #start = time.time()
        #xgb_clf.fit(X_train, y_train)
        #end = time.time()

        #time_elapsed = end - start
        #print(time_elapsed)

[ ]: #y_pred = xgb_clf.predict(X_test)

[ ]: #dtrain = xgb.DMatrix(X_train, label=y_train)
        #dtest = xgb.DMatrix(X_test, label=y_test)

[ ]:
[ ]:
[ ]:

```